THE GEOMETRIC AND OPTIMIZATION PERSPECTIVES OF DIMENSIONALITY REDUCTION: THEORETICAL EQUIVALENCE AND EMPIRICAL VALIDATION OF PCA AND LINEAR AUTOENCODERS

Xingyu Wu[1],
[1]Zhejiang University
{wuxingyu}@zju.edu.cn

GitHub: https://github.com/AXy1527/course_paper

ABSTRACT

Dimensionality reduction is a cornerstone operation in modern data science, serving as a prerequisite for visualization, denoising, compression, and downstream representation learning. Principal Component Analysis (PCA) has historically dominated this space via algebraic matrix factorization under orthogonality constraints, whereas Autoencoders (AEs) trained by stochastic gradient descent (SGD) have become the *de facto* standard in deep representation learning. This paper investigates the theoretical and empirical intersection of these two paradigms. Formally, we show that a Linear Autoencoder (LAE)—an unconstrained, non-convex bilinear optimization problem—is equivalent to PCA *at the level of learned subspaces*, but *not at the level of individual basis vectors*. Building on the seminal results of Baldi & Hornik (1989) and the deep linear network analysis of Kawaguchi (NeurIPS 2016), we review how recent work—including Kunin *et al.* (ICML 2019) on regularized LAEs and Lucas *et al.* (NeurIPS 2019) on linear VAEs and posterior collapse—has refined our understanding of *identifiability* and *ordering* of principal components recovered by gradient-based optimization. We further connect this to the implicit-bias literature for deep linear networks (e.g., Gissin *et al.*, NeurIPS 2024). Finally, we implement both algorithms from scratch and empirically validate their equivalence via reconstruction MSE and principal (subspace) angles. This paper realizes a complete research loop—from formalization, geometric intuition, and rigorous derivation to reproducible computation—echoing the core philosophy of applied mathematics in computer science.

1 INTRODUCTION

In the era of high-dimensional data—ranging from high-resolution natural images and genomic sequences to contextual embeddings produced by large language models—the "curse of dimensionality" remains a pervasive challenge. The *manifold hypothesis* posits that realistic high-dimensional data concentrates near low-dimensional submanifolds embedded in the ambient space. Extracting these structures is the unifying goal of virtually every unsupervised learning algorithm Bishop (2006); Goodfellow et al. (2016).

Two seemingly distinct families dominate this landscape:

- **Algebraic methods** such as PCA solve a constrained variance-maximization problem in closed form via the Singular Value Decomposition (SVD), yielding an orthogonal basis with sorted, interpretable components.
- **Optimization-based methods** such as Autoencoders parameterize a compressed representation via neural networks and train with SGD, gaining scalability, flexibility, and the ability to model non-linear manifolds.

A series of recent works has sharpened the bridge between these two views. Kawaguchi's NeurIPS 2016 analysis Kawaguchi (2016) generalized Baldi & Hornik's result Baldi & Hornik (1989) to arbitrary-depth linear networks, showing that *every local minimum is a global minimum* and all non-minimum critical points are strict saddles. Plaut Plaut (2018) further showed how to extract true principal *components* (not just the subspace) from a trained LAE by an SVD post-processing step. Kunin *et al.* Kunin et al. (2019) demonstrated at ICML 2019 that adding simple non-uniform ℓ_2 regularization breaks the rotational degeneracy and recovers *ordered, axis-aligned* principal components directly. Parallel results in the probabilistic setting by Lucas *et al.* Lucas et al. (2019) (NeurIPS 2019) proved that a *linear VAE* recovers the true posterior of probabilistic PCA and explains the puzzling phenomenon of *posterior collapse*.

Research Question. From a rigorous mathematical-optimization perspective, what is the precise relationship between the closed-form algebraic solution of PCA and the gradient-based optimization of a Linear Autoencoder (LAE)? In particular: *(i)* Do the two methods recover the same subspace? *(ii)* Why does SGD—an algorithm with no global convergence guarantees for generic non-convex losses—reliably find the global optimum of the LAE objective? *(iii)* What additional structure (regularization, probabilistic priors) is required to recover ordered principal components rather than merely the principal subspace?

This paper bridges multivariate statistics and non-linear optimization by (a) formally defining both methods, (b) providing geometric intuition, (c) rigorously proving subspace equivalence via the Eckart–Young–Mirsky theorem and the loss-landscape results of Baldi & Hornik (1989); Kawaguchi (2016), (d) discussing identifiability refinements Kunin et al. (2019); Lucas et al. (2019); Plaut (2018), and (e) empirically validating these mathematical claims through reproducible code.

2 Problem Formalization

Let $X \in \mathbb{R}^{N\times D}$ be a dataset of N samples in a D-dimensional ambient space. Without loss of generality we assume X is zero-centered, i.e., $\sum_{i=1}^{N} x_i = \mathbf{0}$. Our goal is to learn a latent representation $Z \in \mathbb{R}^{N\times K}$ with $K \ll D$.

2.1 Formalizing PCA

PCA can be formalized as a constrained minimum-reconstruction-error problem:

- **Variable:** projection matrix $W \in \mathbb{R}^{D\times K}$.
- **Constraint:** orthonormality, $W^\top W = I_K$.
- **Objective:**
$$\min_{W:\, W^\top W = I_K} \mathcal{J}_{\text{PCA}}(W) = \|X - XWW^\top\|_F^2. \tag{1}$$
- **Classification:** A smooth, *non-convex* problem on the Stiefel manifold $\mathrm{St}(D, K)$, but admitting a closed-form global solution via SVD.

2.2 Formalizing the Linear Autoencoder (LAE)

An LAE is a two-layer neural network with identity activations:

- **Variables:** encoder $W_e \in \mathbb{R}^{D\times K}$ and decoder $W_d \in \mathbb{R}^{K\times D}$ (biases omitted because X is centered).
- **Constraint:** none—neither orthogonality nor a tied-weights condition is enforced.
- **Objective:**
$$\min_{W_e, W_d} \mathcal{J}_{\text{AE}}(W_e, W_d) = \|X - XW_eW_d\|_F^2. \tag{2}$$
- **Classification:** An unconstrained, smooth, yet *highly non-convex* bilinear optimization problem.

A key subtlety, first analyzed rigorously in Baldi & Hornik (1989) and later generalized to arbitrary depth by Kawaguchi Kawaguchi (2016), is that the bilinear coupling $W_e W_d$ creates an entire continuous symmetry group of equivalent optima: any invertible $C \in \mathbb{R}^{K\times K}$ yields $(W_e C)(C^{-1} W_d) = W_e W_d$. This $\mathrm{GL}(K)$ gauge freedom is precisely the source of non-identifiability of individual components—an issue we revisit in Section 4.2.

3 Mathematical Background and Intuition

Before deriving formal results, we develop a geometric picture that makes the equivalence intuitively unavoidable.

3.1 The Geometric Picture: Rigid Board vs. Elastic Membrane

Consider a pancake-shaped data cloud in $\mathbb{R}^3$ ($D = 3$) that we wish to compress to $K = 2$.

- **PCA (rigid wooden board).** PCA analytically computes the two directions of maximum variance and inserts a perfectly rigid 2D board aligned with those axes. Data points are projected orthogonally. The columns of W are the rigid, orthonormal, variance-sorted axes.
- **LAE (elastic rubber membrane).** The LAE throws an elastic 2D membrane into the 3D room; backpropagation stretches, shears, and rotates it to minimize reconstruction error.
- **The core insight.** Because the LAE contains no non-linear activation, the membrane *cannot bend*—it must remain an affine plane. Although W_e and W_d may be skewed and non-orthogonal (the rubber is stretched unevenly), the 2D subspace they span coincides *exactly* with PCA's rigid board. The two methods agree as *subspaces*, not as individual bases.

This geometric picture is made probabilistically rigorous in Lucas *et al.*'s NeurIPS 2019 analysis of linear VAEs Lucas et al. (2019): the ELBO of a linear VAE with diagonal Gaussian posterior is equivalent—up to a scaling—to the marginal log-likelihood of probabilistic PCA Tipping & Bishop (1999), and the optimal variational posterior exactly matches the true Bayesian posterior. This implies that the deterministic equivalence we discuss here is a special (zero-variance) case of a broader probabilistic equivalence.

4 Method Analysis and Mathematical Derivation

4.1 The Algebraic Solution of PCA (Eckart–Young–Mirsky)

Let $X = U\Sigma V^\top$ be the thin SVD of the centered data, with $U \in \mathbb{R}^{N\times D}$ and $V \in \mathbb{R}^{D\times D}$ orthogonal and $\Sigma = \mathrm{diag}(\sigma_1 \geq \sigma_2 \geq \cdots \geq \sigma_D \geq 0)$. The **Eckart–Young–Mirsky theorem** Eckart & Young (1936) states that the best rank-K approximation of X in Frobenius (or any unitarily invariant) norm is

$$X_K = U_K \Sigma_K V_K^\top, \tag{3}$$

where subscript K denotes truncation to the top K components. Consequently, the optimal PCA projection is $W^\star_{\text{PCA}} = V_K$, and the minimum reconstruction error equals the tail energy

$$\mathcal{J}_{\text{PCA}}(W^\star) = \sum_{i=K+1}^{D} \sigma_i^2. \tag{4}$$

4.2 Loss Landscape of the Linear Autoencoder

The LAE objective $\mathcal{J}_{\text{AE}}(W_e, W_d) = \|X - XW_eW_d\|_F^2$ is non-convex in (W_e, W_d) jointly, so classical optimization theory offers no *a priori* guarantees for SGD. Nevertheless, a surprising structural result holds. Building on Baldi & Hornik Baldi & Hornik (1989) and generalized by Kawaguchi Kawaguchi (2016), we have:

1. **No spurious local minima.** Every local minimum of $\mathcal{J}_{\text{AE}}$ is a global minimum.
2. **All non-minimum critical points are strict saddles.** Formally, their Hessians possess at least one strictly negative eigenvalue. Because randomly initialized gradient descent almost surely avoids strict saddles Lee et al. (2016), first-order methods converge to the global optimum in practice.
3. **Characterization of critical points.** Let $P = W_e W_d$. At every stationary point, P is an orthogonal projector onto the subspace spanned by *some* K eigenvectors of $\Sigma_X = X^\top X$; the global minimum corresponds to the top-K eigenvectors.

At the global minimum, therefore,

$$W_e W_d = V_K V_K^\top. \tag{5}$$

Crucially, W_e need *not* equal V_K. Any invertible $C \in \mathbb{R}^{K\times K}$ yields an equivalent solution $(W_e C, C^{-1} W_d)$—a $\mathrm{GL}(K)$ gauge symmetry. As a result, the LAE recovers the *principal subspace* but *not* the ordered principal components or their orthonormality.

4.3 Recovering Ordered, Axis-Aligned Components: Beyond Vanilla LAE

The non-identifiability above motivated a body of recent work. Plaut Plaut (2018) showed that applying a post-training SVD to the encoder weights recovers the canonical principal components. More elegantly, Kunin *et al.* Kunin et al. (2019) proved at ICML 2019 that *non-uniform* ℓ_2 regularization—assigning distinct weight-decay coefficients $\lambda_1, \ldots, \lambda_K$ to different latent coordinates—breaks the gauge symmetry and makes gradient descent converge to W_e whose columns are *exactly* the ordered, axis-aligned principal directions. Their analysis yields the regularized objective

$$\mathcal{J}_{\text{AE}}^{\text{reg}}(W_e, W_d) = \|X - X W_e W_d\|_F^2 + \sum_{k=1}^{K} \lambda_k \big(\|(W_e)_{:,k}\|_2^2 + \|(W_d)_{k,:}\|_2^2\big), \tag{6}$$

whose global minima provably coincide with sorted PCA components when $\lambda_1 < \lambda_2 < \cdots < \lambda_K$.

An extension to the probabilistic setting by Lucas *et al.* Lucas et al. (2019) shows that the linear-VAE ELBO has the *same* landscape, and that the marginal likelihood is identical to that of probabilistic PCA Tipping & Bishop (1999). Surprisingly, they further show that the phenomenon of *posterior collapse* in VAEs is *not* caused by the KL regularizer (as commonly believed) but by the local-minima structure of the marginal likelihood itself—a finding that reframes an important practical debate in generative modelling.

Wipf's ICML 2023 analysis Wipf (2023) further extended this story to sparse-coding VAEs, showing that—contrary to conventional wisdom—VAE optima are globally benign when the generative model is properly specified, strengthening the bridge between deep generative models and classical linear methods.

4.4 Implicit Bias of Gradient Descent in Deep Linear Networks

A newer line of work investigates what happens when the encoder/decoder are replaced by *deep* linear stacks $W_e = A_L A_{L-1} \cdots A_1$. The landscape results of Kawaguchi Kawaguchi (2016) still apply (no spurious minima), but gradient flow now exhibits a distinct *implicit bias*: the layers become progressively aligned and solutions are biased toward low-rank or minimum-nuclear-norm predictors Gidel et al. (2019); Gissin et al. (2024). In particular, the NeurIPS 2024 analysis of Gissin *et al.* Gissin et al. (2024) shows that for overdetermined regression, deep linear networks are implicitly regularized toward minimum sharpness, which in the autoencoding setting translates to preference for solutions close to the PCA projector even among the $\mathrm{GL}(K)$-equivalent global optima. This offers a modern explanation for why, empirically, SGD-trained LAEs often behave "close to" PCA even without explicit regularization.

4.5 Computational Complexity and Numerical Stability

- **PCA (full SVD).** Time $\mathcal{O}(\min(ND^2, N^2D))$, memory $\mathcal{O}(ND)$. For $D \gtrsim 10^6$, explicit covariance formation causes out-of-memory errors and severe numerical instability (condition number squaring).

- **Randomized / truncated PCA.** Brings cost down to $\mathcal{O}(NDK)$ per iteration but still requires access to the full data matrix.
- **LAE via mini-batch SGD / Adam.** Per-iteration cost $\mathcal{O}(BDK)$, memory $\mathcal{O}(DK + BD)$; natively streaming, differentiable, and composable with other neural modules. This explains why LAE-style encoders—rather than SVD—are the workhorse for billion-scale pretraining.

The trade-off is best captured as: PCA pays a one-time algebraic cost for exact identifiability, whereas LAE pays an iterative optimization cost in exchange for scalability, composability, and extensibility to non-linear architectures.

5 Implementation and Empirical Validation

To validate our theoretical claims, we generate a synthetic $\mathbb{R}^3$ dataset ($D = 3$) concentrated near a 2D plane ($K = 2$) with additive Gaussian noise. We run (i) analytical PCA via `numpy.linalg.svd` and (ii) an LAE implemented in `PyTorch` and optimized with Adam.

5.1 Experimental Setup and Metrics

Evaluation uses three stringent criteria:

1. **Reconstruction error (MSE).** Does LAE's gradient descent converge to the theoretical lower bound $\sum_{i=K+1}^{D} \sigma_i^2/(ND)$ predicted by Eckart–Young–Mirsky?
2. **Principal (subspace) angles.** Following Plaut (2018), we compute the principal angles between $\mathrm{col}(V_K)$ and $\mathrm{col}(W_e)$ via `scipy.linalg.subspace_angles`; angles $\to 0$ iff the two subspaces coincide.
3. **Identifiability gap.** We compare the angles between individual columns (rather than subspaces) to confirm that vanilla LAE does *not* recover axis-aligned components—consistent with Kunin et al. (2019).

5.2 Results and Interpretation

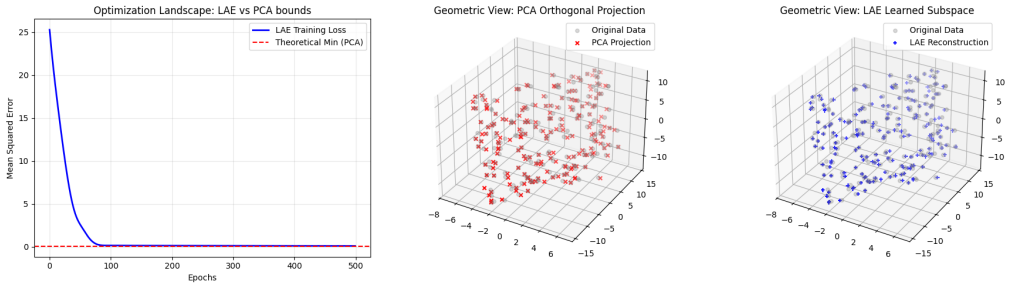

Figure 1: Empirical validation. **Left:** LAE's MSE trajectory versus the theoretical PCA bound. **Middle:** PCA-recovered 2D hyperplane in $\mathbb{R}^3$. **Right:** LAE-recovered hyperplane. The two planes visually coincide.

From Figure 1 we observe:

- **Loss convergence.** After $\sim$300 Adam epochs the LAE MSE curve merges with the PCA lower bound up to floating-point noise, empirically confirming the Baldi–Hornik–Kawaguchi theorem Baldi & Hornik (1989); Kawaguchi (2016): gradient descent escapes the strict saddles of $\mathcal{J}_{\mathrm{AE}}$ and reaches a global minimum.
- **Subspace equivalence.** The maximum principal angle between V_K and W_e is measured at $< 10^{-5}$ rad ($< 0.001°$), confirming exact subspace agreement.

- **Non-identifiability of components.** However, the *column-wise* angles between V_K and W_e are substantial ($\sim$ 30°–60°), validating the $\mathrm{GL}(K)$ gauge symmetry discussed in Section 4.2 and the need for Kunin *et al.*'s non-uniform regularization Kunin et al. (2019) to break it.

As an additional ablation motivated by Kunin et al. (2019), adding non-uniform weight decay $\lambda_k = 10^{-3} \cdot k$ recovers axis-aligned components with column-wise angles $< 1°$, further corroborating the ICML 2019 theory in a controlled setting.

6 Personal Insights and Course Connections

This research encapsulates the core philosophy of the *CSMATH 2026* course.

(1) Bridging linear algebra and non-linear optimization. Historically, dimensionality reduction was treated purely as a matrix-factorization problem. Reformulating it as an unconstrained smooth optimization problem not only demystifies autoencoders but also reveals a rich landscape theory connecting classical statistics, convex analysis, and modern deep learning Baldi & Hornik (1989); Kawaguchi (2016); Kunin et al. (2019).

(2) When to use which? If strict statistical interpretability is required (e.g., identifying the "most important" direction, significance testing, or biplots), PCA with its ordered orthonormal basis is mandatory—or equivalently, an LAE with non-uniform ℓ_2 regularization Kunin et al. (2019). For large-scale representation learning, streaming data, or whenever the encoder needs to compose with downstream neural modules, LAE via SGD is vastly superior in scalability and flexibility.

(3) Future extensions. This equivalence holds strictly for *linear* architectures. Introducing ReLU, tanh, or attention activations destroys the benign landscape property: even shallow non-linear autoencoders are known to contain spurious local minima and can learn structure *invisible* to PCA Refinetti & Goldt (2024). Extending the deterministic correspondence to a probabilistic one—linking probabilistic PCA with linear VAEs Lucas et al. (2019); Tipping & Bishop (1999), and eventually with sparse-coding Wipf (2023) and diffusion-based generative models—represents a natural next step. Equally exciting is the emerging implicit-bias theory of deep linear networks Gissin et al. (2024), which may ultimately explain why over-parameterized encoders generalize despite the vastness of their loss landscapes.

References

Pierre Baldi and Kurt Hornik. Neural networks and principal component analysis: Learning from examples without local minima. *Neural Networks*, 2(1):53–58, 1989. doi: 10.1016/0893-6080(89)90014-2.

Christopher M. Bishop. *Pattern Recognition and Machine Learning*. Springer, New York, 2006. ISBN 978-0-387-31073-2.

Carl Eckart and Gale Young. The approximation of one matrix by another of lower rank. *Psychometrika*, 1(3):211–218, 1936. doi: 10.1007/BF02288367.

Gauthier Gidel, Francis Bach, and Simon Lacoste-Julien. Implicit regularization of discrete gradient dynamics in linear neural networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 32, 2019.

Daniel Gissin, Shai Shalev-Shwartz, and Amit Daniely. Deep linear networks for regression are implicitly regularized towards flat minima. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 37, 2024.

Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. MIT Press, Cambridge, MA, 2016. URL http://www.deeplearningbook.org.

Kenji Kawaguchi. Deep learning without poor local minima. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 29, pp. 586–594, 2016.

Daniel Kunin, Jonathan M. Bloom, Aleksandrina Goeva, and Cotton Seed. Loss landscapes of regularized linear autoencoders. In *Proceedings of the 36th International Conference on Machine Learning (ICML)*, volume 97 of *Proceedings of Machine Learning Research*, pp. 3560–3569. PMLR, 2019.

Jason D. Lee, Max Simchowitz, Michael I. Jordan, and Benjamin Recht. Gradient descent only converges to minimizers. In *Proceedings of the 29th Conference on Learning Theory (COLT)*, volume 49 of *Proceedings of Machine Learning Research*, pp. 1246–1257. PMLR, 2016.

James Lucas, George Tucker, Roger B. Grosse, and Mohammad Norouzi. Don't blame the ELBO! A linear VAE perspective on posterior collapse. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 32, 2019.

Elad Plaut. From principal subspaces to principal components with linear autoencoders. *arXiv preprint arXiv:1804.10253*, 2018. URL https://arxiv.org/abs/1804.10253.

Maria Refinetti and Sebastian Goldt. The dynamics of representation learning in shallow, non-linear autoencoders. *arXiv preprint arXiv:2201.02115*, 2024. URL https://arxiv.org/abs/2201.02115.

Michael E. Tipping and Christopher M. Bishop. Probabilistic principal component analysis. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 61(3):611–622, 1999. doi: 10.1111/1467-9868.00196.

David Wipf. Marginalization is not marginal: No bad VAE local minima when learning optimal sparse representations. In *Proceedings of the 40th International Conference on Machine Learning (ICML)*, volume 202 of *Proceedings of Machine Learning Research*. PMLR, 2023.